

Software Transactional Memory and Microservices: A Match Made in Heaven

Mark Little, VP Red Hat UK Ltd

Michael Musgrove, Red Hat UK Ltd


17th Sept, 2019

Agenda

- The Actor Model
- Distributed Transactions and Software Transactional Memory (STM)
- Microservices
- Why do these all belong together?
- Technology (Quarkus, Vert.x, Narayana STM)
- Coding demo

The Actor Based Programming Model

- Actors and CSP have been around for decades
 - CSP from Hoare, 1985
 - Actor model from Hewitt et al, 1973
- But popular ways to model primitives for concurrent computations
 - Embodies Processing, Storage and Communication
- Distributed computations communicate via message passing




Nice Features of the Actor Model

- Fixed message sets (i.e., no hidden or unexpected interactions)
- Simplified data management (state is internal to the actor)
- Location transparency (because other actors only see the address)
- Loose coupling
- Asynchronous message passing
 - Has much in common with the reactive manifesto (Responsive, Elastic, Asynchronous and, with extensions, Resilience).

Distributed Transactions

- ACID properties
- Two-phase commit
 - Required when there is more than one resource (RM)
 - Managed by the transaction manager (TM)
 - Uses a familiar two-phase technique (2PC)





Software Transactional Memory

- Software Transactional Memory (STM) proposed in 1995
- STM is about ease of use and reliability
 - Access shared state, either for reading or writing, occurs within atomic blocks
 - All code inside an atomic block executes as if it were single threaded
 - Less error prone (the atomic block is the protection) than traditional concurrency primitives or placing composite operations behind an API
 - Some implementations can be lock free (optimistic vs pessimistic, timestamp)
 - Has some of the same characteristics of ACID transactions

Transactions and Actors

- An actor may go through multiple state transitions upon receipt of a message
- An actor could be internally implemented using multiple threads
- Computational failures may occur
- Hardware and software failures may occur
- Consistency of state important
- Composition of actors
- The combination of STM and Actors is fairly natural

Microservices 101





Nizar S.

@natewave

 **Follow**



Finally, thanks to microservices, my dream of being a detective has come true. Every bug is more like a murder mystery.

RETWEETS

4

LIKES

6



1:18 PM - 11 Jun 2016



4



6



Tweet

**stacks machine** @cemerick · 05/01/2015 ▾

Uh, microservices. So, people are hooking minute bits of computation together via unmanaged pipes carrying opaque chunks of encoded data?



9



36



44



Christian Posta Retweeted

**stacks machine**

@cemerick ▾

Replying to @cemerick

Microservices, because designing, implementing, deploying, monitoring, managing, and supporting network APIs is so fucking easy.

05/01/2015, 20:40

112 RETWEETS 109 LIKES

Tweet your reply




OK so WHAT are they then?!

- Microservices pushes us (back) to distributed systems (circa 1996)
 - Services inherently in separate machines (physical or virtual)
 - Communication between distributed systems slower than within same address space
 - Failures can happen independently
 - Cascading failures can happen
 - Consistency concerns become even more of a challenge
 - Sh*t happens!
 - Throwing developers into distributed systems 101 is not a good way to be agile!

Enterprise microservices

- Microservices distributed systems present challenges
 - The need for transactions, reliable messaging etc. doesn't go away
 - If anything it becomes more important to developers
 - Application containers breaking into pieces
 - Independently deployable (Linux container based) services
 - Available to different language clients using REST/HTTP
 - And other protocols?
 - Still A LOT slower than in-memory!

Notable Technology Used in the Demonstration



OPENSIFT



VERT.X

@narayana

Quarkus




A cohesive platform for optimized developer joy:

- Based on standards, but not limited
- Unified configuration
- Zero config, live reload in the blink of an eye
- Streamlined code for the 80% common usages, flexible for the 20%
- No hassle native executable generation

WAIT.
SO YOU JUST SAVE IT,
AND YOUR CODE IS RUNNING?
AND IT'S JAVA?!



I KNOW, RIGHT?
SUPERSONIC JAVA, FTW!




Boot time to compile time



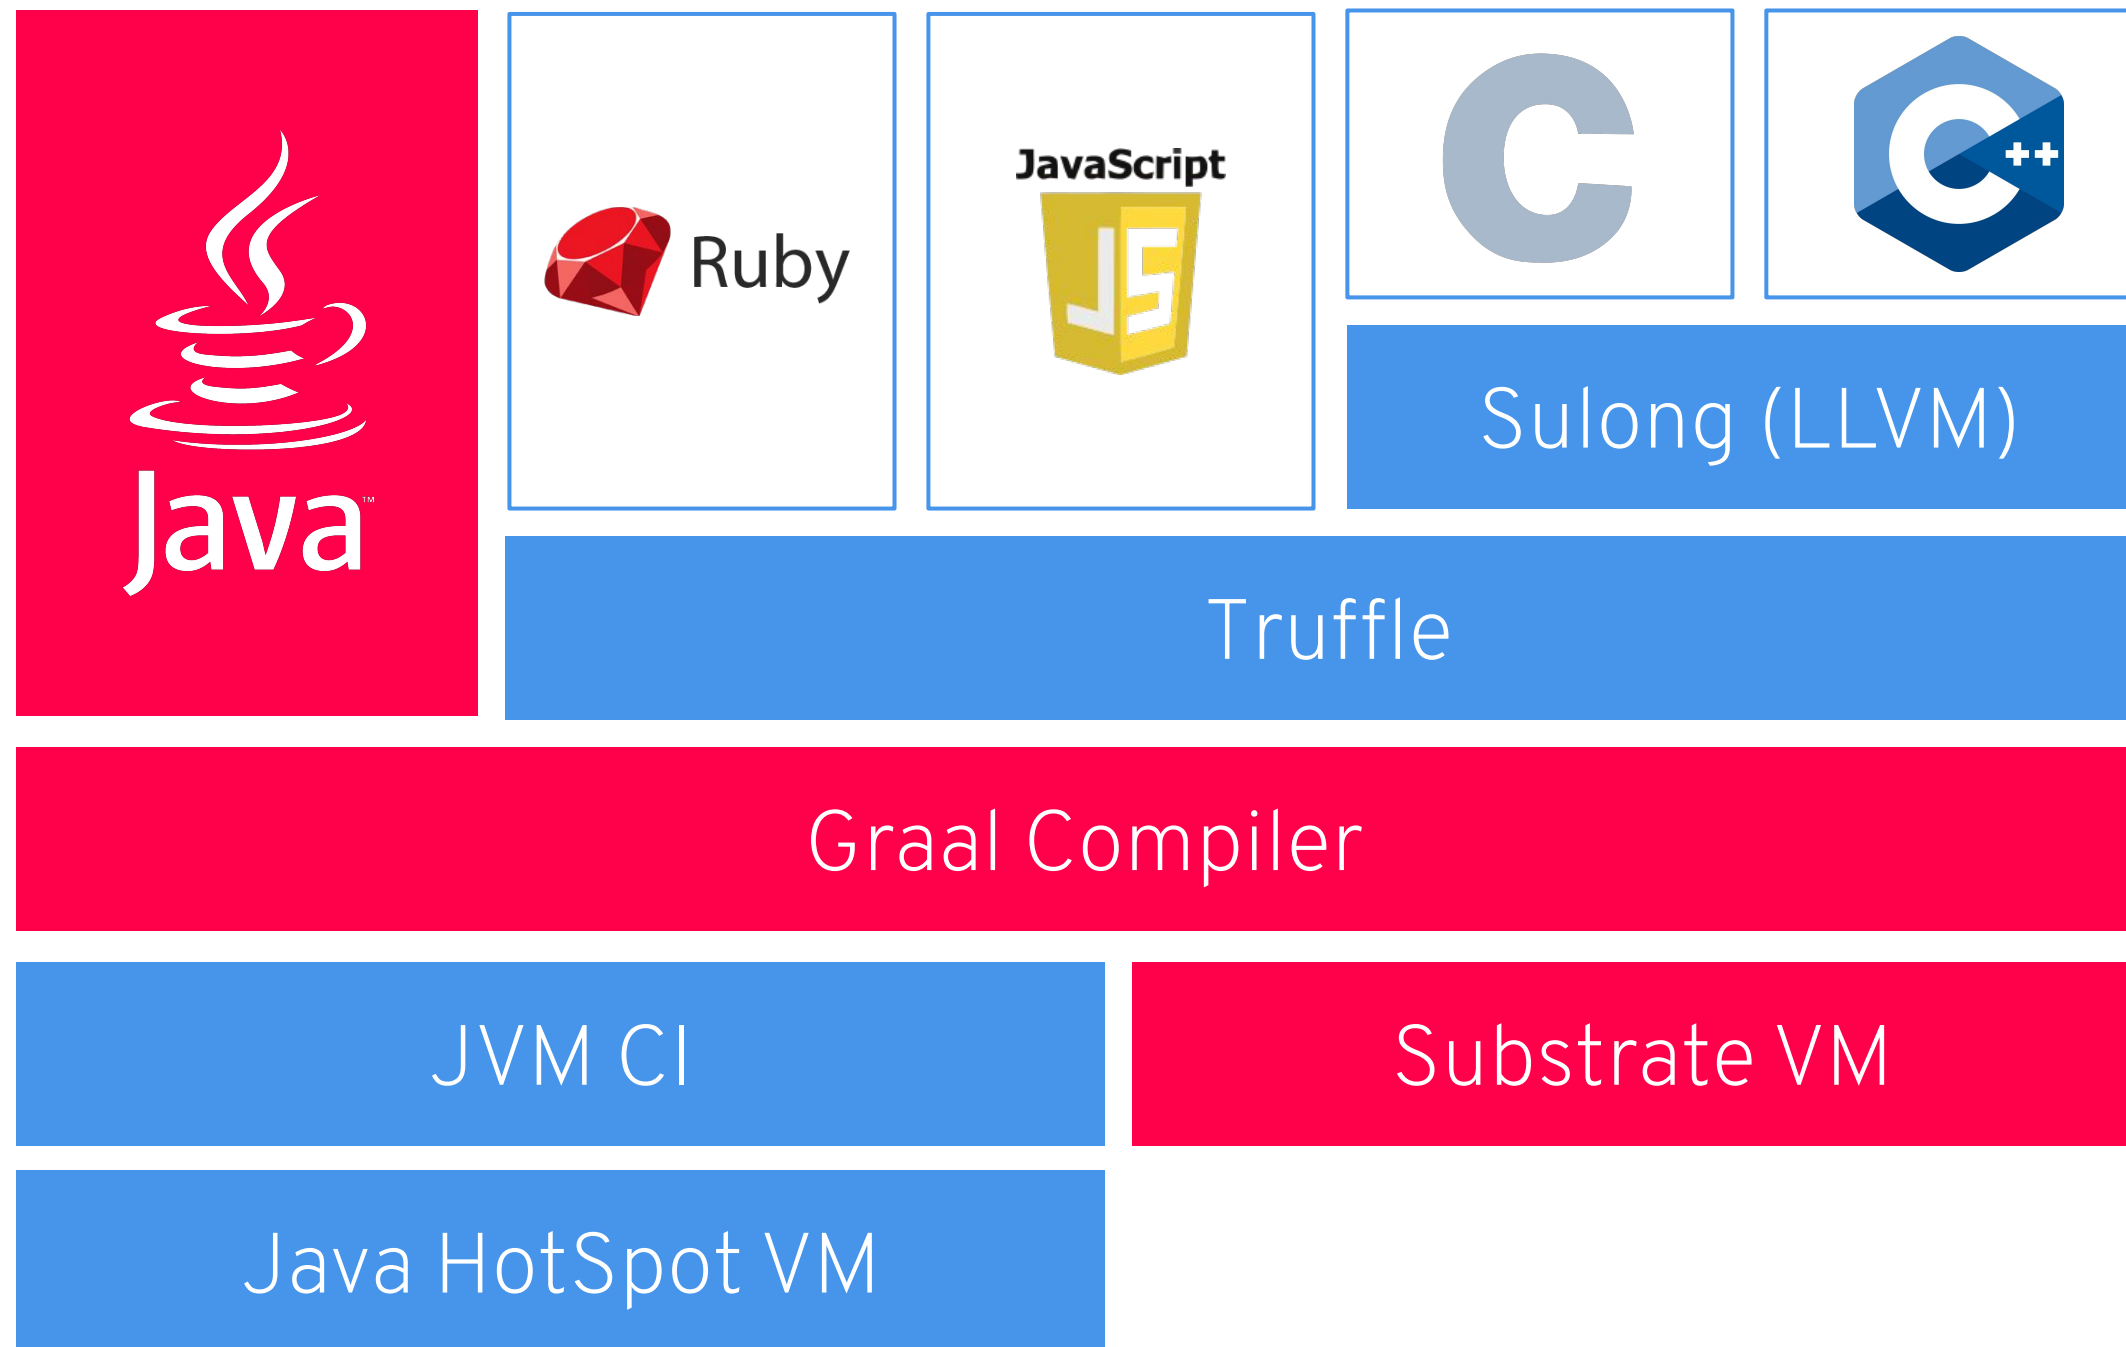
What does a framework do at startup time?

- Parse config files
- Classpath & classes scanning
 - for annotations, getters or other metadata
- Build framework metamodel objects
- Prepare reflection and build proxies
- *Start and open IO, threads etc*






Framework Optimizations

- Moved as much as possible to build phase
- Minimized runtime dependencies
- Maximize dead code elimination
- Introduced clear metadata contracts
- Spectrum of optimization levels
(all → some → no runtime reflection)




Supersonic Subatomic Java

REST




Supersonic Subatomic Java

REST



REST + CRUD



Time to first response

Eclipse Vert.x

- Non-blocking (asynchronous) and non locking (concurrency is natural)
- Event bus (reactor pattern)
- Load balancing, failover, circuit breaker
- Clustering and Service Discovery
- Polyglot JVM
- Infinispan/JDG and Spring Boot
- Added AMQ and Qpid Dispatch Router
- TCP, UDP, HTTP 1 & 2 servers and clients (non-blocking), gRPC





Narayana STM

- Provides an STM implementation
- Define state (objects) which can be manipulated within transactions
 - Volatile (recoverable) and persistent (durable)
- Pessimistic and Optimistic models
- Different variants of transactions
 - Top level
 - Nested
 - Nested top level
- Modularity

STM Object Annotations

- `@Transactional`
 - Implementations of the interface then become container managed
- `@Nested` & `@NestedTopLevel`
 - The container will create a new transaction for each method
- `@Optimistic` & `@Pessimistic` (with `@Timeout` and `@Retry` options)
- `@ReadLock` & `@WriteLock`
- `@State` & `@NotState`
- `@TransactionFree`

OpenShift



Microsoft Azure



Code demonstration

An STM example running on Quarkus

Three versions of an application running on Quarkus:

1. A standard async JAX-RS app
 - a. No concurrency control
 - b. Reactive/async support via Vert.x
 - c. Stress test to show concurrency issues
2. Same app using volatile STM for concurrency control
 - a. Make the service interface a volatile STM object and restress
3. Deploy onto Openshift
 - a. Turn it into a shared persistent STM object and deploy to OpenShift
 - b. Create a shared persistent volume for the STM logs
 - c. Test and then scale up and down via the OpenShift console to show persistency

Where can I find out more?

- More information available from <https://narayana.io>:
 - Forums
 - Blogs
 - IRC
- Demo source
 - <https://github.com/narayana/conferences/tree/master/codeone-2019>
- STM source
 - <https://github.com/jbosstm/narayana/tree/master/STM>

THANK YOU

- any questions?



plus.google.com/+RedHat



facebook.com/redhatinc



linkedin.com/company/red-hat



twitter.com/RedHatNews



youtube.com/user/RedHatVideos